

## 1.) Import the data from CCLE into a new Google Colab file

```
import pandas as pd
from google.colab import drive
import matplotlib.pyplot as plt
import statsmodels.api as smf
```

```
from google.colab import drive
drive.mount('/content/gdrive/', force_remount=True)
```

Mounted at /content/gdrive/

```
import sklearn as sk
from sklearn.linear_model import LinearRegression
import statsmodels.api as sm
```

```
insurance=pd.read_csv("/content/gdrive/MyDrive/Econ441B/insurance.csv")
insurance.head()
```

	age	sex	bmi	children	smoker	region	charges
0	19	female	27.900	0	yes	southwest	16884.92400
1	18	male	33.770	1	no	southeast	1725.55230
2	28	male	33.000	3	no	southeast	4449.46200
3	33	male	22.705	0	no	northwest	21984.47061
4	32	male	28.880	0	no	northwest	3866.85520

```
insurance.loc[insurance['sex']=='female','sex']=1
insurance.loc[insurance['sex']=='male','sex']=0
insurance.loc[insurance['smoker']=='yes','smoker']=1
insurance.loc[insurance['smoker']=='no','smoker']=0
```

```
from sklearn import preprocessing
insurance['region']=preprocessing.LabelEncoder().fit_transform(insurance['region'])
insurance.head()
```

	age	sex	bmi	children	smoker	region	charges
0	19	1	27.900	0	1	3	16884.92400
1	18	0	33.770	1	0	2	1725.55230
2	28	0	33.000	3	0	2	4449.46200
3	33	0	22.705	0	0	1	21984.47061
4	32	0	28.880	0	0	1	3866.85520

## 2.) Split the data into 80/20, in/out sample

```
cut=int((len(insurance)*0.8//1))
```

```
in_data=insurance.iloc[:cut,0:6]
out_data=insurance.iloc[cut:,0:6]
in_target=insurance.iloc[:cut,6]
out_target=insurance.iloc[cut:,6]
```

## 3.) Normalize the Data

```
from sklearn import preprocessing
scaler=preprocessing.StandardScaler().fit(in_data)
in_data_scale = scaler.transform(in_data)
out_data_scale = scaler.transform(out_data)
```

#### 4.) Get lambda from Lasso cross validation

```
from sklearn.linear_model import LassoCV
modCV=LassoCV().fit(in_data_scale, in_target)
a=modCV.alpha_
a

133.34880015958146
```

#### 5.) Run a lambda regression with that Lambda

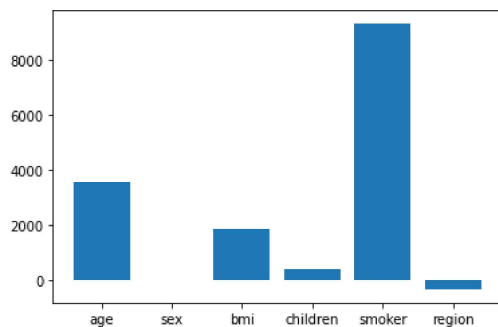
```
from sklearn.linear_model import Lasso
modl=Lasso(alpha=a).fit(in_data_scale, in_target)

modl.coef_

array([3564.26137833,  0.          , 1890.8527485 ,  398.24050447,
       9324.84227138, -326.52198757])
```

#### 6.) Visualize the coefficients

```
plt.bar(['age', 'sex', 'bmi', 'children', 'smoker', 'region'], modl.coef_)
plt.show()
```



#### 7.) Interpret the coefficients

From the plot in part6, we can see that, age, bmi, children, smoker is positively correlated with response variable charges. And region (southwest=3, southeast=2, northwest=1, northeast=0) is negatively correlated with response variable charges. And sex can be removed from the model.

#### 8.) Compare in and out of sample MSE's

```
import numpy as np
from sklearn.metrics import mean_squared_error
from math import sqrt
predict_in_data=modl.predict(in_data_scale)
print(mean_squared_error(in_target, predict_in_data))
predict_out_data=modl.predict(out_data_scale)
print(mean_squared_error(out_target, predict_out_data))

36490415.101693384
37252730.724018715
```

As we can see above, the MSE for in-sample data is 36490415.10, which is a little bit lower than the out-sample data which is 37252730.72

✓ 0s completed at 3:17 PM

● ×